

R for Data Analysis in Agrobiodiversity

1

Intro and data types



Introduction to the course

- The course is organized in 7 lessons (+exam)
- 20 hours total (2 CFU)
- Attendance to 70% of the lectures is required
- Theory + practice; we will start from the basics, and aim to data analysis. Let's see how it goes

Planned programme

Program of the course:

L1: Introduction to R, installation, data structures

Aula 1 (Palazzo Toscanelli), 03/04/2019, 9.30 - 12.30

L2: Basic scripting, file management

Aula 1 (Palazzo Toscanelli), 04/04/2019, 14.30 - 17.30

L3: Plotting, output to figures and text

Aula 1 (Palazzo Toscanelli), 05/04/2019, 14.30 - 17.30

L4: Advanced scripting, exploratory data analysis

Aula 10 (Main Building), 11/04/2019, 13.30 - 16.30

L5: Data analysis I

Sala Riunioni, 2° Floor (Palazzo Toscanelli), 15/04/2019, 9.30 – 12.30

L6: Data analysis II


Aula 10 (Main Building), 18/04/2019, 13.30 - 16.30

L7: Data analysis III, wrap-up

TBD

Some instructions

- Always bring your laptop, fully charged
- The program is not written in stone; let me know what you are looking for and we will try to arrange it
- Be flexible; many different levels of preparation in the room
- You will learn much more R by exercising on your own than sitting in this room
- Don't be afraid if you loose something; scripts are there for this reason, and will be shared
- Google is your best friend; learn to ask the right questions



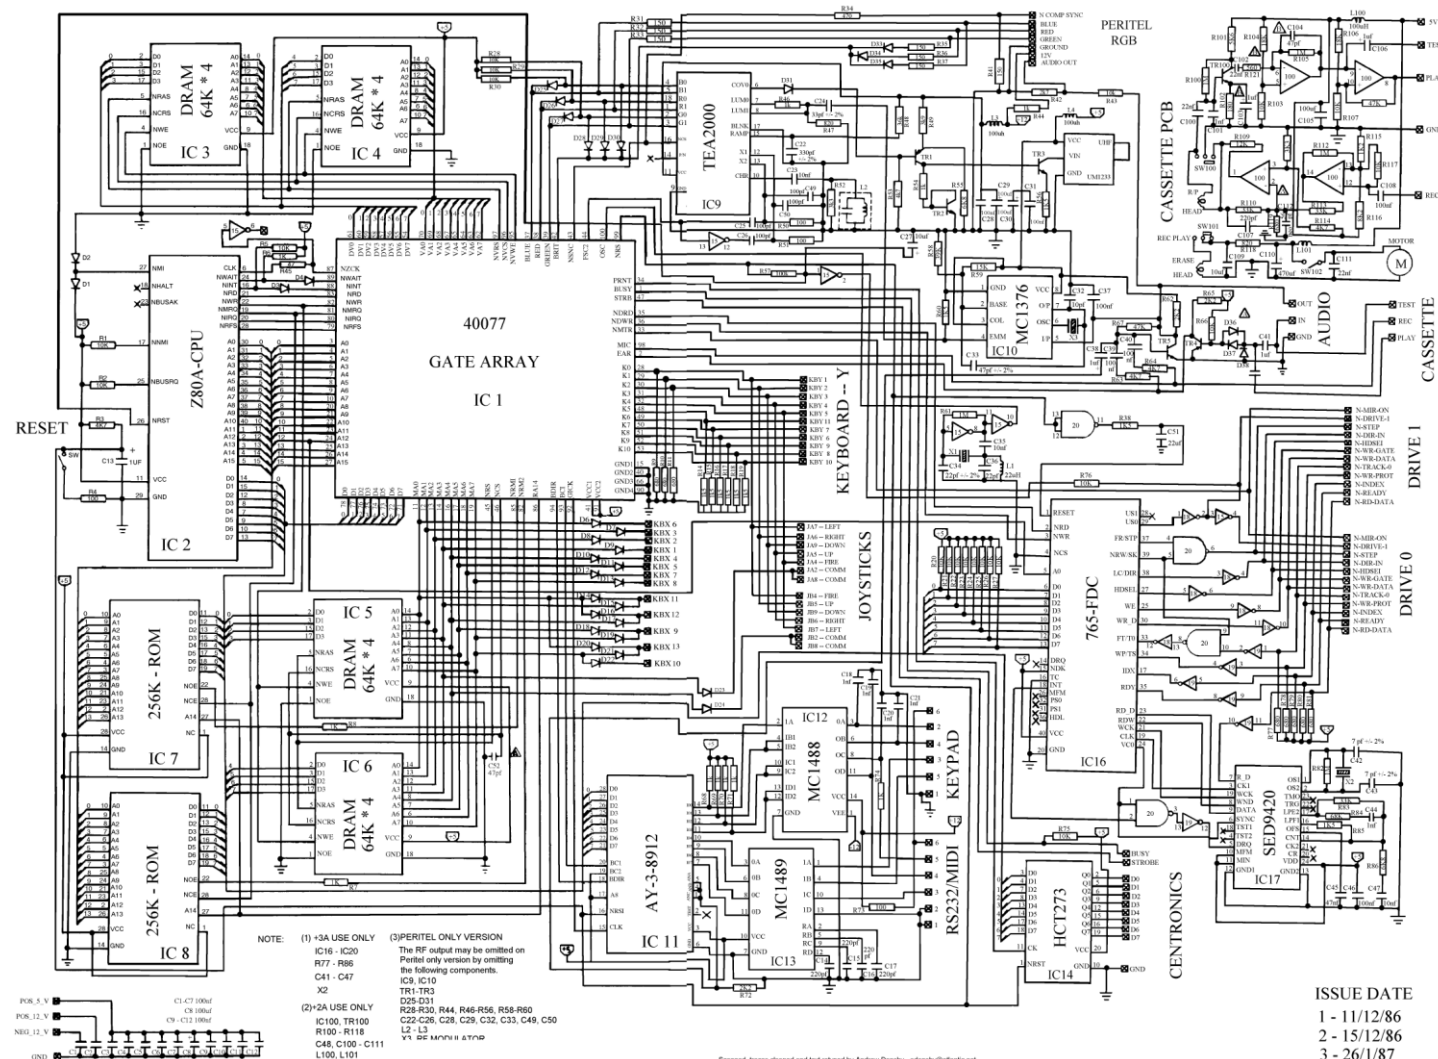
Who am I (to teach you R)?

- NOT an informatician
- Assistant Professor in crop genetics here at the Scuola
- R user for some time now, mainly using R in genomics and genetic data analysis
- Pro Google user

What I use R for

- Data manipulation, quality check
- Exploratory data analysis
- Statistical analyses
- Plotting for papers and presentations

A few notes on computers



ISSUE DATE
1 - 11/12/86
2 - 15/12/86
3 - 26/1/87

Hardware & Software

- Computer hardware are physical devices used in or with your machine
- Computer software is a collection of code run on the hardware
- SW&HW are interconnected; hardware is useless without software, software cannot be run without hardware

HW

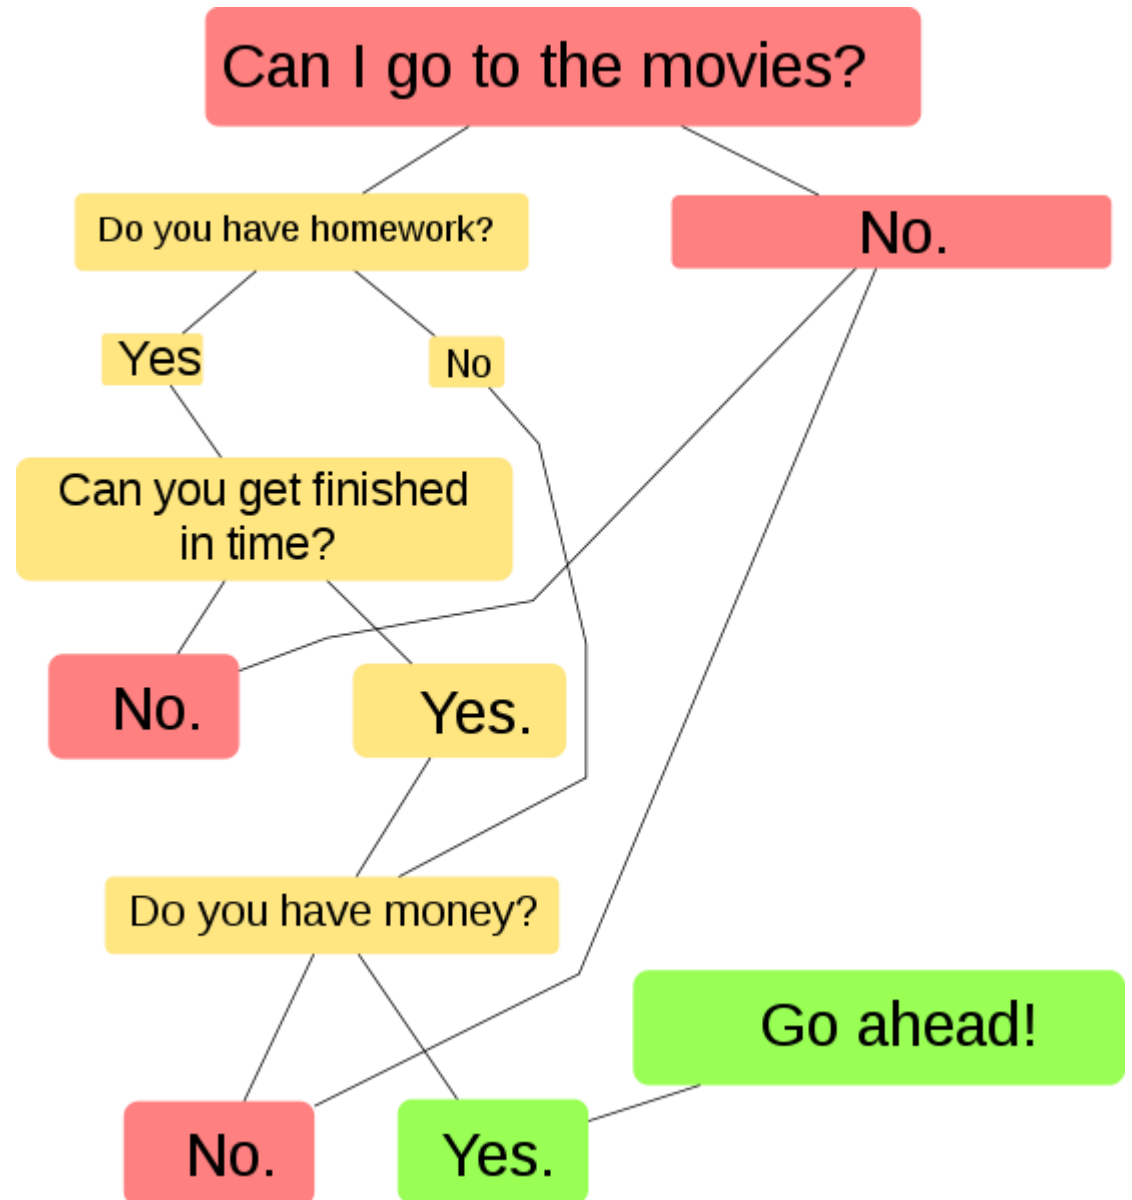


SW



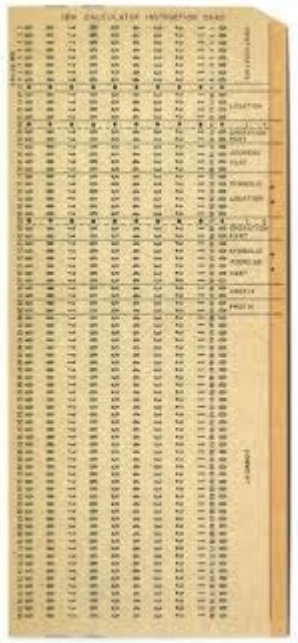
Algorithms

- Computers are dumb; they can't do anything without being told what to do
- Algorithms are step-by-step instructions given to computers in order to solve problems and/or perform tasks



Input devices

- Hardware devices that you use to feed information to the computer



Hard disk drive (HDD)

- Stores all the data necessary to run processes
- It can be magnetic (HDD) or solid state (SSD)
- It is the same found in pen drives



Its capacity determines the **amount of data** that you can store

Think of it as your cabinet. The bigger it is, the more stuff you can put into it




Random access memory (RAM)

- Temporarily stores data, allowing you to work on it
- Provides quick read and write access to the storage device

Its size determines the **amount of data** you can work on **simultaneously**

Think of it as your desk; the bigger it is, the more data you can simultaneously work on



Central processing unit (CPU)

- Comes in many flavours (GPUs are also a thing)
- Responsible for processing information
- Works by logic gates: either there is electric current (1) or there isn't (0)
- It compiles any information in binary code (0/1)
- Might be installed in parallel to increase computing power

The CPU controls the **amount** of simultaneous processes that you can run and their **speed**



Output devices


- Devices (peripherals) that receive data from the computer and send it to the outer world



Motherboard

- The backbone of your computer
- Hosts all main components and puts them in relation
- Host circuits supplying power and circulating information





SOFTWARE

HARDWARE

CPU

OUTPUT


INPUT

«Grandma, I want tortelli!»

HDD


RAM






A few words on programming languages

- Akin to idioms from all over the world, programming languages follow their own grammar and syntax rules
- Different languages share some degree of compatibility depending on their relatedness
- Languages cannot be easily translated in one another



Programming languages are developed to communicate with computers

Computers and humans speak (and **think**) in very different ways



Programming languages can be sorted by their closeness to human language in:

- High level (close to human)
- Low level (close to machine)

Eventually, every instruction must be provided in 0s and 1s



High level language

- Clarity
- Fruibility

Java

R


Interpreter making a translation

C

Assembly

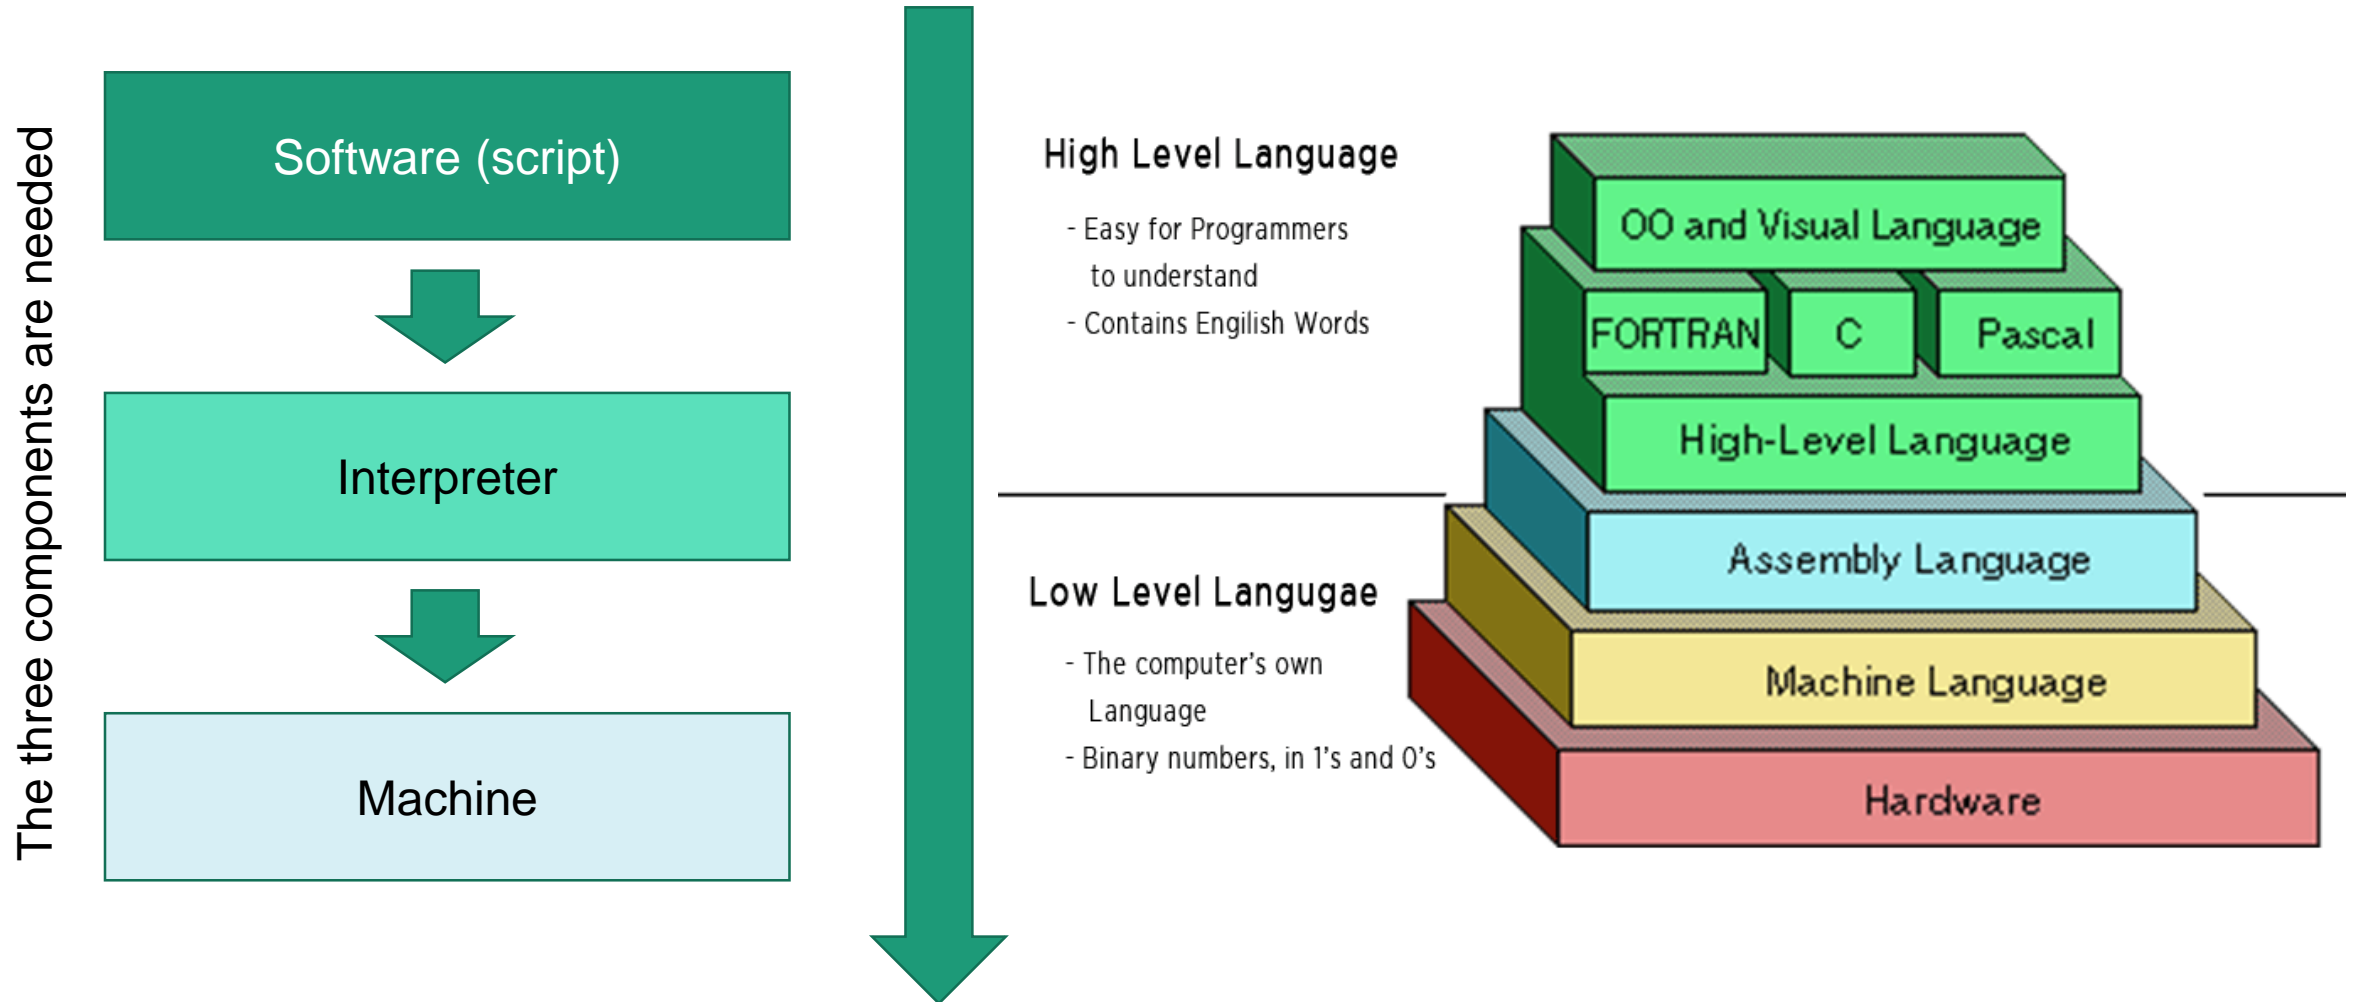
Machine language

- Speed
- Efficiency



Low level language

Computer languages come with interpreters




A friendly reminder



Why are we learning R?

- R is a programming language and an environment for statistical computing
- It is a free software contributed by a large community
- It is one among many programming languages



- It is NOT the fastest
- It is NOT the most efficient
- It is NOT the easiest

However, R is perhaps the language most used by the scientific community (at least in life sciences). For good reasons:

- It is great to conduct statistical analyses
- It is great to produce plots
- It is supported by a large community providing a number of tools aiding data analysis and supporting reproducibility

R installation

Go to R website and download the installation package compatible with your platform (Win, OS, Linux)

<https://cran.r-project.org>

The Comprehensive R Archive Network

Download and Install R

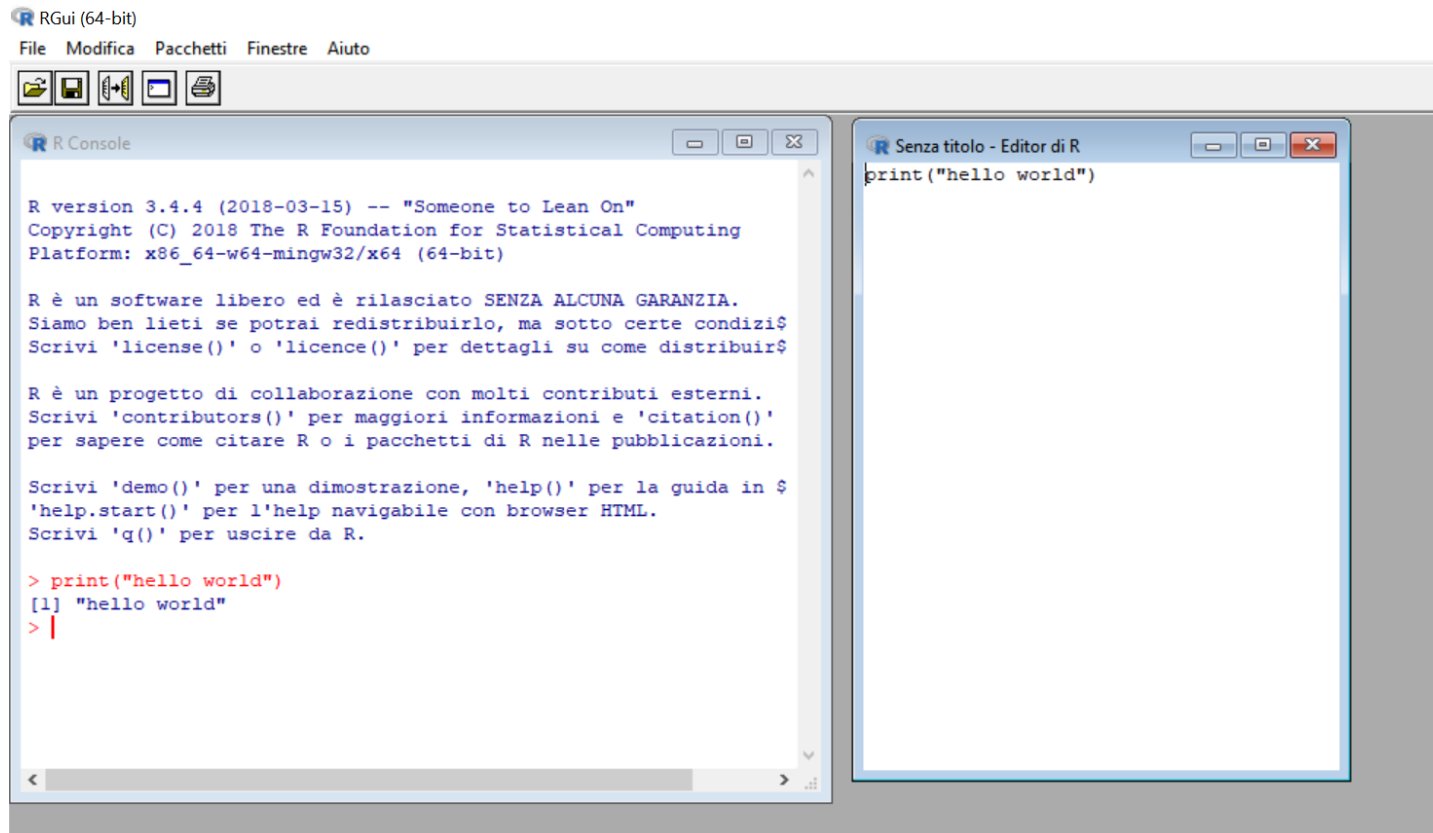
Precompiled binary distributions of the base system and contributed packages, **Windows and Mac** users most likely want one of these versions of R:

- [Download R for Linux](#)
- [Download R for \(Mac\) OS X](#)
- [Download R for Windows](#)

R is part of many Linux distributions, you should check with your Linux package management system in addition to the link above.



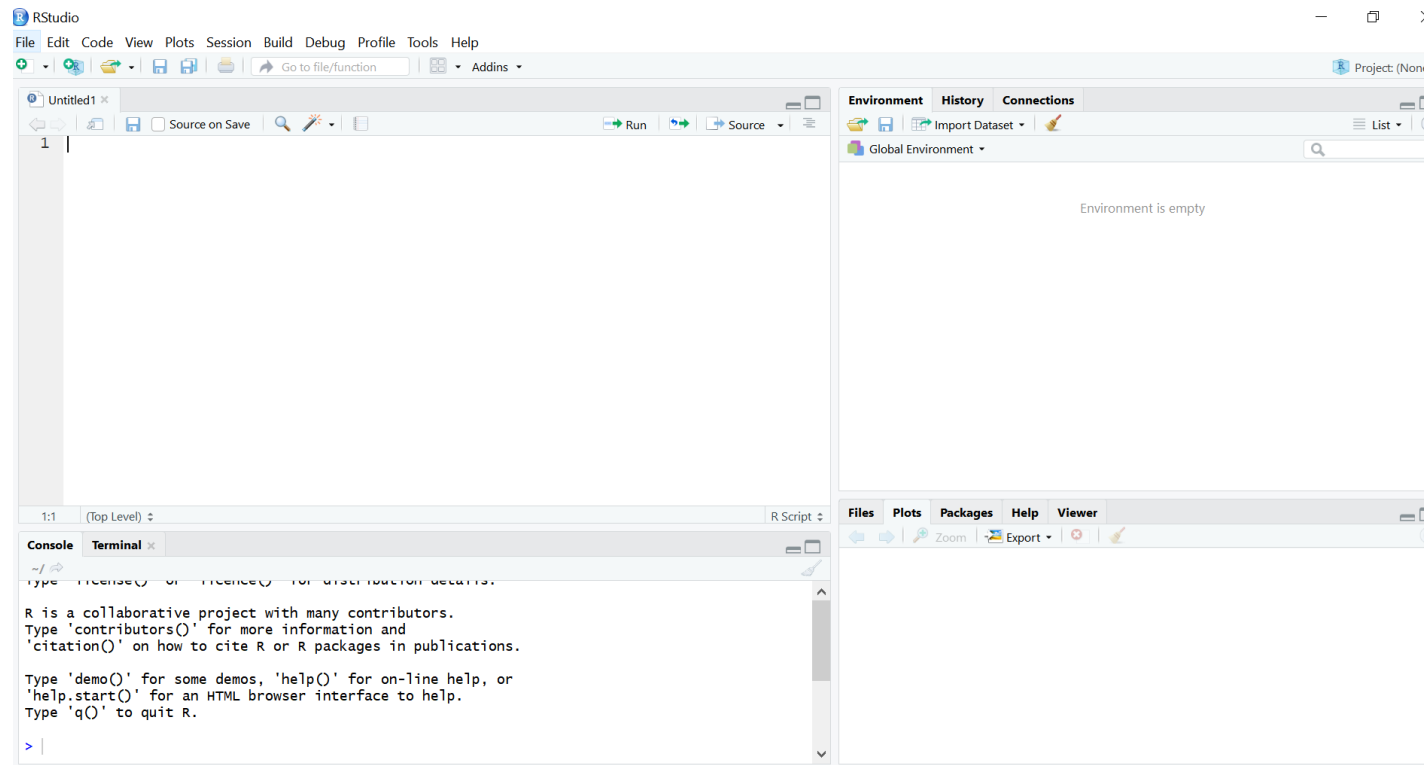
R comes in a basic interpreter



- Not really user friendly
- Not really fancy
- Quite flexible
- Cross-platform
- **Good for learning**

I suggest installing also an advanced text editor such as **Notepad++** (Win) or **SublimeText** (OS)

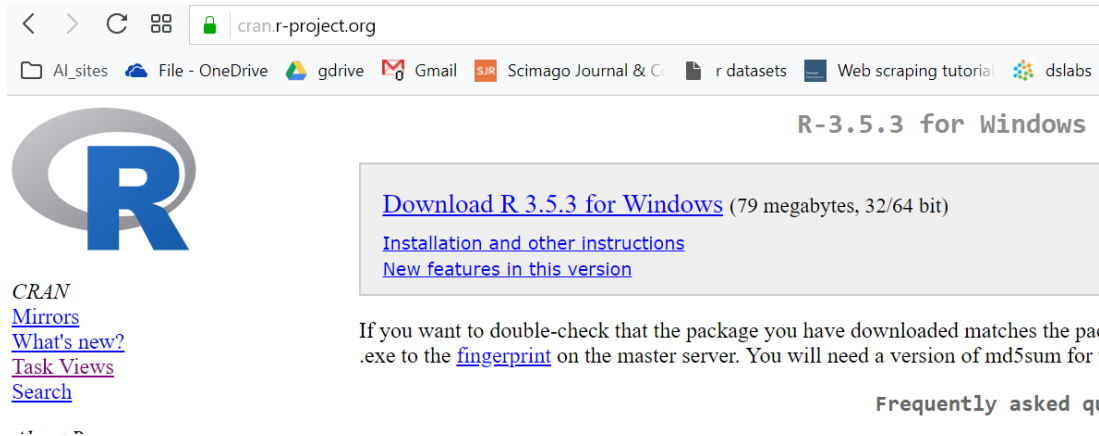
For the sakes of this course, we will use Rstudio, as it provides additional functionalities



Note: Rstudio requires the installation of the R interpreter to work

Move on to install Rstudio and an advanced text editor

<https://cran.r-project.org>



The screenshot shows the CRAN website for R 3.5.3 for Windows. The page features the R logo on the left and a large blue 'R' in the center. Below the logo, there are links for 'CRAN Mirrors', 'What's new?', 'Task Views', and 'Search'. The main content area has a heading 'R-3.5.3 for Windows' and a button to 'Download R 3.5.3 for Windows (79 megabytes, 32/64 bit)'. Below this, there are links for 'Installation and other instructions' and 'New features in this version'. A paragraph of text explains how to verify the download using the 'fingerprint' on the master server. At the bottom, there is a section for 'Frequently asked questions'.

<https://www.rstudio.com>

Choose Your Version of RStudio

RStudio is a set of integrated tools designed to help you be more productive with R. It includes a console, syntax-highlighting editor that supports direct code execution, and a variety of robust tools for plotting, viewing history, debugging and managing your workspace. [Learn More about RStudio features.](#)



RStudio Desktop
Open Source License

RStudio Desktop
Commercial License

RStudio Server
Open Source License

RStudio Server Pro
Commercial License

RStudio Server Pro +
RStudio Connect
Commercial License

FREE

\$995 per year

FREE

\$9,995 per year

\$29,995 per
year

DOWNLOAD

[Learn More](#)

BUY

[Learn More](#)

DOWNLOAD

[Learn More](#)

DOWNLOAD

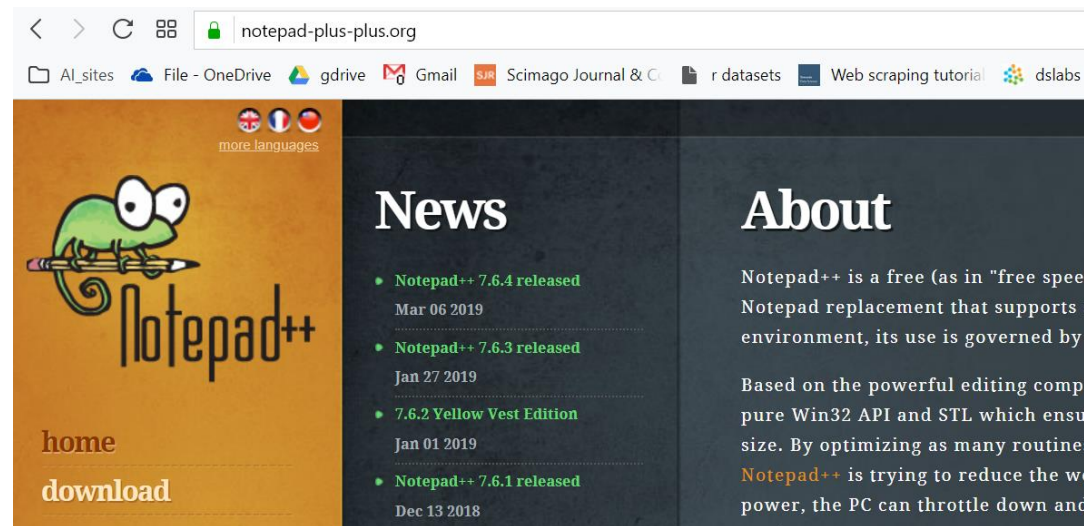
[Learn More](#)

TALK

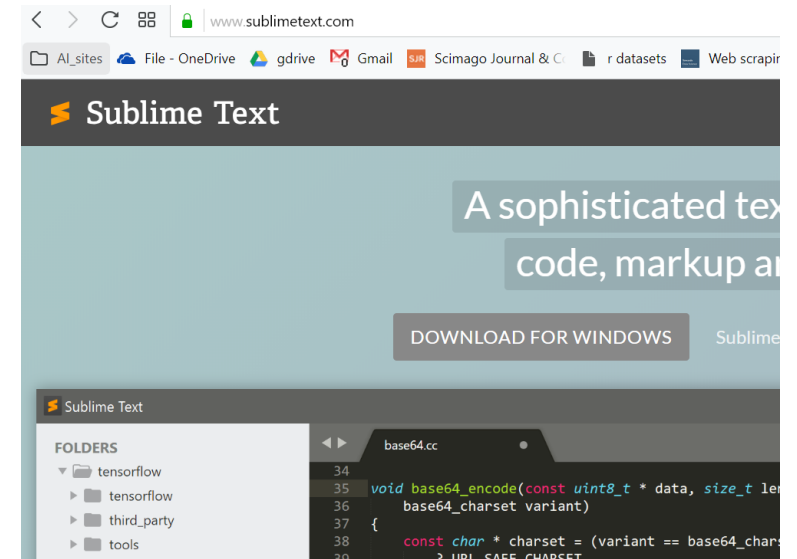
[Learn More](#)

<https://www.sublimetext.org>

<https://www.notepad-plus-plus.org>



The screenshot shows the Notepad++ website. The page has a dark background with a green frog logo on the left. The main content area is divided into three columns: 'News', 'About', and 'Download'. The 'News' column lists several releases of Notepad++ with dates. The 'About' column contains text about the software being free and its features. The 'Download' column has a large button to download the software. At the bottom, there are links for 'home' and 'download'.



The screenshot shows the Sublime Text website. The page has a dark background with the Sublime Text logo at the top. The main content area features a large heading 'A sophisticated text editor' and a button to 'DOWNLOAD FOR WINDOWS'. Below this, there is a preview of the Sublime Text editor interface showing a code editor with a file named 'base64.cc' and a sidebar with a file tree.

Let's get to coding

When working in an R console, “>” is the command prompt; you will find it at the beginning of each line of code, provided that the console is ready to accept commands

The hashtag character “#” is used to denote comments; any line of text after it is not considered code

Each line of code is stand-alone, provided that special characters are not used to extend it in subsequent lines

Lower case letters and uppercase letter are considered different: Hello is not hello

You may use how many white spaces as you want in between operators

R is a powerful calculator

```
> 2*3  
[1] 6
```

```
> 7*8  
[1] 56
```

```
> 5+7*8  
[1] 61
```

```
> (5+7)*8  
[1] 96
```

```
> 4^2  
[1] 16
```

```
> 100%%10  
[1] 0
```

```
> 100%%11  
[1] 1
```

Arithmetic operators

Operator	Description
+	Addition
-	Subtraction
*	Multiplication
/	Division
^	Exponent
%%	Modulus (Remainder from division)
%/%	Integer Division

Be ALWAYS aware of operator precedence

```
> 3 > 2  
[1] TRUE
```

```
> 3 != 3  
[1] FALSE
```

```
> 3*2 == 2*3  
[1] TRUE
```

```
> 3^2 <= 2^3  
[1] FALSE
```

Relational operators

Operator	Description
<	Less than
>	Greater than
<=	Less than or equal to
>=	Greater than or equal to
==	Equal to
!=	Not equal to

Functions

- The cool thing is that you can assemble basic operators in what are called functions
- This helps in wrapping simple procedures in more complex tools, that can be nested one within the other (we will see more about that)
- Functions are identified by the round bracket operators

Function	Description
<code>sqrt()</code>	root square
<code>log()</code> , <code>log10()</code>	logarithm
<code>exp(a)</code>	e^a
<code>mean(a)</code>	mean of vector a
<code>min(a)</code> , <code>max(a)</code>	min/max values of a
<code>sd(a)</code>	standard deviation of a


`function.name(argument1, argument2, ...)`

R is object-oriented programming (OOP)

Values are assigned to objects through assignment operators

```
> 3*5  
[1] 15  
  
> a <- 3*5  
  
> a  
[1] 15
```

Operator	Description
<-, <<-, =	Leftwards assignment
->, ->>	Rightwards assignment



Operators and functions may then be used on objects

```
> a <- 3+3
```

```
> b <- 3*3
```

```
> a==b  
[1] FALSE
```

```
> c<- a*b
```



```
> c < b  
[1] FALSE
```

As a matter of fact, we will use objects all the time

Object names are **case-sensitive**:

“Result” is different from “result”

There are a few reserved words that is better not to use to name objects, especially those used by functions (e.g. “mean”)



- Depending on their features, objects may be traced back to classes
- Classes are characterized by specific attributes
- Classes define the data format (and features) of objects

R objects may belong to a number of different classes

«integer»

«logical»

«character»

«numeric»

«factor»

[...]

Classes dictate what you can do on a specific object; you cannot multiply characters, you cannot search for strings in logicals, etc.

Watch out for the representation of values

Integer

3

Floating number

3.0

Character

"3"

Coercion of data types

- Data may be converted from one type to the other with the appropriate functions: `as.numeric()`, `as.character()`, ...

```
> a<-"3"  
> anum<as.numeric(a)  
> anum  
[1] 3
```

- Be EXTRA careful and double check that everything is as intended

```
> is.numeric(anum)  
[1] TRUE
```

Data structures

Data loaded in R may assume different structures, depending on their dimensionality and features

- Vectors
- Matrices
- Dataframes
- Lists
- Arrays

Atomic vectors

- Ordered collections of objects of the same type
- Represent one-dimensional data
- May be created using the `c()` function that concatenates all its arguments

```
> vector <- c(1,4,7,2)
> vector
[1] 1 4 7 2
> class(vector)
[1] "numeric"
```

```
> vector <- c(1,4,7,"2")
> vector
[1] "1" "4" "7" "2"
> class(vector)
[1] "character"
```

Be careful with the dynamic assignment of object classes

Some functions for vectors

`length()`

number of elements in vector

`sum()`

sum of all elements of vector

`cumsum()`

cumulative sum

`prod()`

product of all elements of vector

`order()`

order vector elements

`head()`

print the beginning of the vector

`unique()`

extract unique values

Try it

Make a vector named "vec" with the following elements:

21, 23, 24, 21, 55, 32

- Calculate length, maximum, product, sum, cumulative sum
- Get unique values in a new vector named "u.vec"

Matrices

- Rectangular set of elements with n rows and m columns
- It contains data of the same type
- Can be initialized with the matrix() function

```
> mat <- matrix("value", nrow=4, ncol=5, byrow=T)
> mat
      [,1] [,2] [,3] [,4] [,5]
[1,] "value" "value" "value" "value" "value"
[2,] "value" "value" "value" "value" "value"
[3,] "value" "value" "value" "value" "value"
[4,] "value" "value" "value" "value" "value"
```

Some functions for matrices

`nrow()` / `ncol()`

number of rows/columns in matrix

`colSums()` / `rowSums`

sum by columns/rows

`dim()`

give dimensions of matrix

`rbind()` / `cbind()`

adds rows or columns to matrix

`head()`

print the first few rows of the matrix

`t()`

transpose matrix

`colnames()` / `rownames ()`

get (and set) names of columns and rows

Try it

Make a matrix "mat" with values from 1 to 25 sorted in 5 columns, 5 rows

- Get dimensions, calculate sums by columns
- Create a new object "tmat" transposing mat and calculate sums by rows
- Join "tmat" and "mat" by rows in a new "mat2" object, then compute again sums by columns

Data frames

- Similar to matrices, but able to contain different classes of data at once
- You may initialize them with the `data.frame()` function

```
> vector1<-c("A", "A", "B", "B")
> vector2<-c(32,43,67,65)
> df<-data.frame(type=vector1, count=vector2)
> df
```

	type	count
1	A	32
2	A	43
3	B	67
4	B	65

Try it

Make a df joining the vector "u.vec" (see above) and a vector with the elements "cat", "cat", "dog", "cat", "dog"

- Get number of rows, columns
- Create an object "bigdf" joining "df" and "mat" by rows
- fix row names using a vector with elements "r1", "r2", "r3", "r4", "r5"
- fix column names using numbers from 1 to 7

Lists

- Lists are ordered sets of elements of different classes
- They can contain anything, including other lists
- You may initialize lists with the `list()` function

```
> ls<-list("A", c(1,2,5), list("nested", "list"))
```



```
> ls
[[1]]
[1] "A"

[[2]]
[1] 1 2 5

[[3]]
[[3]][[1]]
[1]
"nested"

[[3]][[2]]
[1] "list"
```

Try it

Make a list "weirdlist" joining mat, mat2, bigdf, and vec

- Get length of list
- Print head of list elements and inspect them

Arrays

- Objects capable to store data in any number of dimensions
- Data stored in arrays must be homogenous
- Arrays might be initialized with the function `array()`

```
>vector1 <- c(5,9,3)
>vector2 <- c(10,11,12,13,14,15)
>arr <- array(c(vector1,vector2),dim = c(3,3,2))
>arr
```

 [,1] [,2] [,3]

[1,]	5	10	13
[2,]	9	11	14
[3,]	3	12	15

 [,1] [,2] [,3]

[1,]	5	10	13
[2,]	9	11	14
[3,]	3	12	15

In short...

	Homogeneous	Heterogeneous
1d	Atomic vector	List
2d	Matrix	Data frame
nd	Array	

Structure of objects at a glance

If you feel lost, you may know the features of your objects with a couple of very useful functions

- `class()`, to give you the class of your object
- `str()`, to give you the structure of your object

```
> df
  type count
1  A    32
2  A    43
3  B    67
4  B    65
```

```
> str(df)
'data.frame':  4 obs. of  2 variables:
 $ type : Factor w/ 2 levels "A","B": 1 1 2 2
 $ count: num  32 43 67 65
```

Data indexing

- To access individual elements of data structures, we may use the square brackets operators: []
- Depending on dimensionality, brackets may be used to slice different components of data using the comma separator «,», the colon operator «:», the dollar symbol «\$», the dash symbol «-»

1. Numeric indexing
2. Name indexing
3. Logical indexing

Indexing vectors

vector[index]

```
> vec<-c(2,5,7,8,9,12,14)
```

```
> vec[5]  
[1] 9
```

Single element

```
>vec[2:5]  
[1] 5 7 8 9
```

Range

```
>vec[c(2,5)]  
[1] 5 9  
> vec[-c(2,5)]  
[1] 2 7 8 12 14
```

Specific
elements

```
> vec[vec>5]  
[1] 7 8 9 12 14
```

Logical
statement

Indexing Matrices/Data Frames

dataframe[row index, col index]

```
> df
  var count
1  A    50
2  B    44
3  C    53
```

```
> df[2,]
  var count
2  B    44
```

By row

```
> df[,2]
[1] 50 44 53
```

By column

```
> df$var
[1] A B C
Levels: A B C
```

By name

```
> df[3,2]
[1] 53
```

Specific
element

```
> df[df[,2]>=50,]
  var count
1  A    50
3  C    53
```

Nested
logical


```
>ls<-list("A", df, c(1:6))
```

```
> ls
```

```
[[1]]
```

```
[1] "A"
```

```
[[2]]
```

```
var count
```

```
1  A   50
```

```
2  B   44
```

```
3  C   53
```

```
[[3]]
```

```
[1] 1 2 3 4 5 6
```

Indexing Lists

list[[index]]

```
> ls[[1]]
```

```
[1] "A"
```

Single element

```
> ls[[2]][1:3,2]
```

```
[1] 50 44 53
```

Nested subset

Indexing Arrays

array[dimension 1, dimension 2, dimension n]

Indexing may also be used in negatives

```
> vec  
[1] 2 5 7 8 9 12 14  
  
> vec[-c(1:3)]  
[1] 8 9 12 14
```

All of the above is **VERY** useful to create new objects and manipulate data

```
> df  
  var count  
1  A    50  
2  B    44  
3  C    53  
  
> df[-1,]  
  var count  
2  B    44  
3  C    53
```

Handling missing data

- Missing data may be stated in any object with the special value **NA** (not available)
- **NA** values are treated differently from other values and may be independently targeted

```
> x<-c(1,5,NA,3)
> is.na(x)
[1] FALSE FALSE  TRUE FALSE
```

- Other special values are **Inf** (infinity) and **NaN** (Not a Number)

Be careful, as some functions and processes will need special instructions to deal with missing data!!

In whichever software you are coding your data, there should not be any other missing data value

- ?

- *

- Missing

- 0

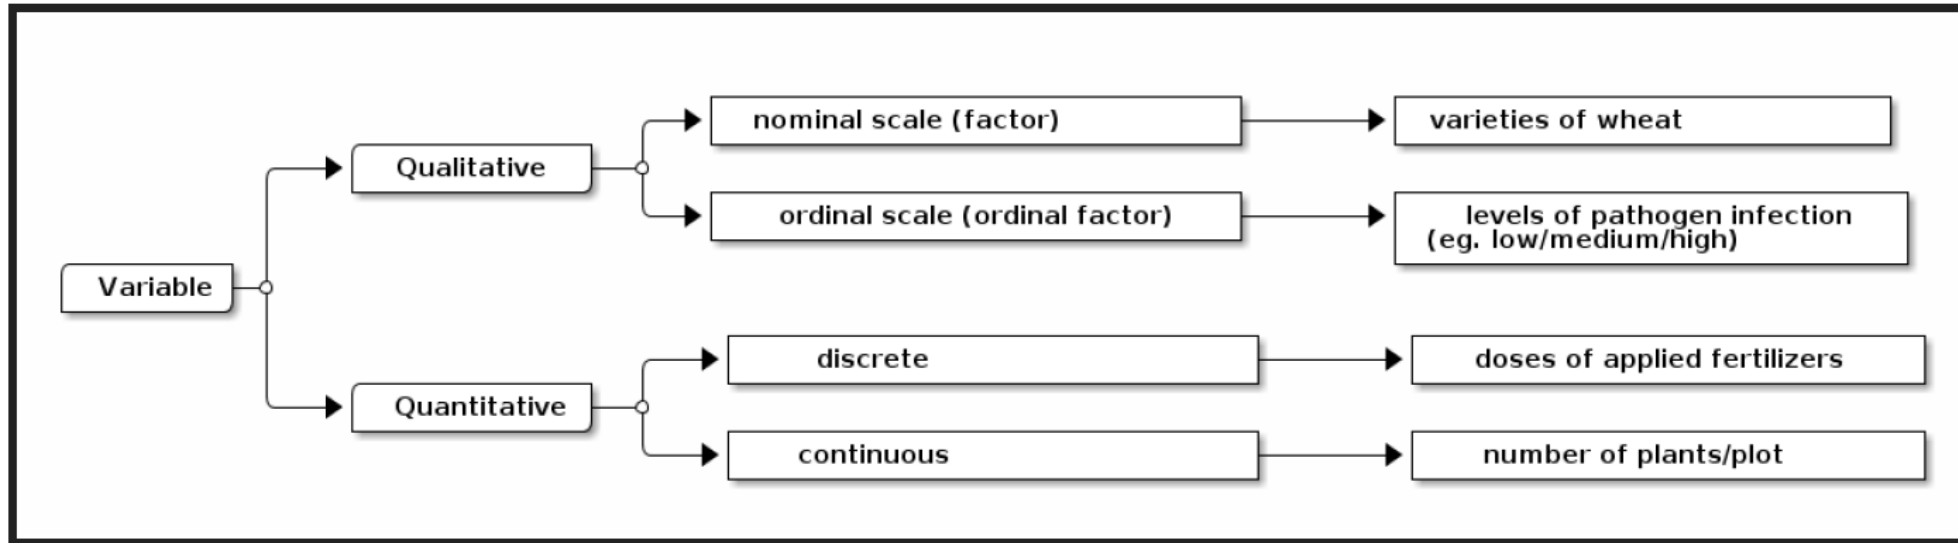


Especially bad

- ...

These will affect the class of the data imported by R, most notably converting numbers in strings or factors

In statistics, variables come in many flavours



They are each best represented by an object class

The factor class: handle with care

- Variables that take only a limited number of different values
- Good for representing categorical variables and conducting the appropriate statistics (e.g. ANOVA)
- Factors are stored as vectors with a corresponding set of values (factors) printed when the vector is printed

```
> vec<-c(1,2,3,1,2,3,3,3,3,2,1,2,4)
> vec
[1] 1 2 3 1 2 3 3 3 3 2 1 2 4
> fvec<-factor(vec)
> fvec
[1] 1 2 3 1 2 3 3 3 3 2 1 2 4
Levels: 1 2 3 4
```

Watch out, as levels are embedded in the object and will remain there also after filtering

```
> fvec  
[1] 1 2 3 1 2 3 3 3 3 2 1 2 4  
Levels: 1 2 3 4  
> fvec[13]<-2  
> fvec  
[1] 1 2 3 1 2 3 3 3 3 2 1 2 2  
Levels: 1 2 3 4
```

In any moment, you may know which are the levels included in a factor object with the function *levels(object)*

If you are not planning to use factors, for good use it may be good to set *option(stringsAsFactors=F)* at the beginning of your script; otherwise, character strings may be interpreted as factor levels

Recap exercise

Load some flower information by typing `data(iris)`



Rookies:

1. What type is this data?
2. What class is data in column 1? Column 5?
3. How many flowers have a width greater than 1?
4. Create a new matrix with sepal length, sepal width, and species name

Pros:

1. What is the mean petal length in the versicolor species?

Recap exercise 2

Load data from an experiment testing the relation between vitamin C and tooth growth in guinea pigs --> `data(ToothGrowth)`

Rookies:

1. Get the structure of the dataset
2. Create two datasets based on unique values in column 2
3. Get dimensions of each
4. What is the mean tooth length in each dataset?
5. What is the mean tooth length in each dataset but only for doses of 2.0?



Pros:

1. Create as many datasets as there are supp, and doses in each; compute mean tooth length in each, and assemble them in a new dataframe with 2 rows (each for supp) and three columns (each for a dose)